

Beyond aggregating a building index: spatially honest, event-conditioned pluvial flood learning on adaptive H3 grids

Working manuscript (methods / IJDRR-shaped)

Status: Matured draft constrained to live Lower Manhattan open-data pilot metrics. Incomplete items marked 待补充.

Public code and documentation: <https://github.com/Coucou2016/pluvial-flood-risk-DGGS-H3>

Terminology ledger

Term	Canonical meaning
H3	Uber hexagonal Discrete Global Grid System
Open labels	DEP stormwater polygons, 311 flooding points, USGS Ida HWM — not PF1b
Spatial H3-block CV	GroupKFold holding out coarse H3 parent blocks
$PFI_h(c, r)$	Model rainfall-conditioned flood probability/index for cell c under rainfall condition r
PF1b	7Analytics / Svellingén building-level pluvial flood index — not used here
LM pilot	Lower Manhattan open-data pilot ($n_cells=141$)

Abstract

Urban pluvial flood screening increasingly relies on machine learning and multi-resolution grids, yet many operational indices remain tied to proprietary insurance labels and evaluation protocols that ignore spatial leakage. Building on the H3 DGGS framing popularised for scalable pluvial mapping, we present an **open-label hexagonal learning protocol** for a Lower Manhattan pilot: multi-source public flood indicators are joined to H3 cells, gradient-boosting models are assessed with **H3-block spatial cross-validation**, scale-loss is quantified with a Jaccard/F1 hotspot ladder, and an **adaptive** refinement stage screens parents using trained cell scores. We define $PFI_h(c, r)$ explicitly as a rainfall-conditioned hex flood probability/index — a model output, not feature importance and not the proprietary PF1b product. On the live open-data pilot table ($n=141$ cells; open-data assembly), spatial CV accuracy was 0.784 ± 0.069 (5 folds) with mean F1 **0.866**; continuous-risk R^2 under the same folds was near zero (0.030 ± 0.343) and is not claimed as skill. Adaptive refinement used about **57%** as many cells as a uniform R11 grid while refining high-score parents — a **cell-count** comparison only. We do **not** claim citywide skill, PF1b reproduction, radar rainfall, or rainfall-conditioned

discrimination: event rainfall remains a synthetic constant hook, and scenario tables show **no within-cell PFI_h change across rainfall intensities** (待补充). Oslo synthetic runs remain appendix-only pipeline QA.

1 Introduction

Intense short-duration rainfall can overwhelm urban drainage and generate pluvial flooding with limited warning. City-scale assessment needs representations that are (i) computationally scalable, (ii) updateable when new observations arrive, and (iii) evaluable without silent spatial leakage. Discrete Global Grid Systems (DGGS), and in particular Uber H3, provide nested hexagonal cells that support multi-resolution aggregation and neighbour queries.

Svelling et al. (2026) demonstrated that machine-learning-derived building-level pluvial susceptibility (PFIb) can be aggregated into H3 cells to reduce spatial query cost by roughly two orders of magnitude, while also showing that hotspot sets can diverge sharply across resolutions (Jaccard similarity ≈ 0.14 between fine street-level and coarser neighbourhood hotspots on their proprietary stack). That work motivates H3 as a communication and screening fabric, but it does not provide an open, spatially honest learning protocol for jurisdictions that cannot access insurance claims.

This manuscript therefore asks: **can an open multi-source label stack on H3 support spatially blocked evaluation, diagnose scale-loss, and drive adaptive refinement with an explicit rainfall-conditioned cell index — without claiming PFIb reproduction?** We contribute a **methods protocol**, not a citywide operational product:

1. An open-label H3 feature/label assembly protocol with explicit provenance fields for assembly mode, feature source, label source, and rainfall source.
2. Primary skill reporting via spatial H3-block CV, with random splits demoted to diagnostics.
3. A scale-loss Jaccard/F1 ladder and an adaptive H3 ablation screened by trained PFI_h, plus a binding definition of $\text{PFI}_h(c, r)$ that is neither feature importance nor PFIb.

Evidence is restricted to a Lower Manhattan open-data pilot ($n=141$). Claims of citywide skill, radar rainfall, PFIb parity, or rainfall-conditioned discrimination are out of scope until those gaps close (待补充).

2 Related work

From building indices to hexagonal screening. Escalating short-duration extremes have pushed operational risk assessment toward leaner, updateable representations. Svelling et al. (2026, IJDRR) show that a machine-learning building-level pluvial susceptibility index (PFIb) can be aggregated into H3 cells, cutting geometry-heavy spatial query cost by roughly two orders of magnitude ($\sim 98\%$ in their report) while exposing a clear resolution trade-off: fine street-level hotspots are largely invisible at coarser neighbourhood grids (Jaccard ≈ 0.14 between their R13 and R10 hotspot sets). A related SSRN preprint develops the same H3

indexing narrative (Svelling et al., SSRN). That line of work establishes H3 as a scalable communication fabric for proprietary indices; it does **not** supply an open-label learning protocol or spatially blocked evaluation for jurisdictions without insurance claims. We therefore treat PFIb→H3 aggregation as the closest precedent and as an explicit claim boundary (no PFIb reproduction; no numerical equality to their Jaccard).

Hexagonal DGGS as analysis substrate. Independently of insurance labels, Li et al. (2022) use an ISEA3H hexagonal DGGS for multi-scale flood mapping under climate scenarios, emphasising resolution-consistent predictors. Their contribution supports DGGS nesting as a scientific substrate. Our focus differs: an Uber H3 **learning and evaluation protocol** on open urban flood indicators, not climate-scenario inundation mapping alone.

Spatial holdouts in GeoAI. Random train/test splits routinely inflate skill under spatial autocorrelation. Sun, Hu, Lakhanpal & Zhou (2023) summarise block / GroupKFold spatial cross-validation for GeoAI; flood susceptibility studies increasingly adopt analogous holdouts. We instantiate blocking with coarse H3 parent IDs so that entire parent blocks are withheld — aligning primary reporting with leakage-aware practice rather than treating random accuracy as the headline metric.

Urban pluvial ML susceptibility. City-scale studies map susceptibility from DEM, land cover, and drainage proxies with classical ML (e.g. Bersabe & Jun, 2025, Seoul). Few combine open multi-source labels, H3 nesting, adaptive refinement screened by trained scores, and an explicit rainfall-conditioned cell index in one reproducible pipeline with blocked CV as the primary metric. That combination — and the honesty constraints that follow from a small open-data pilot — is the gap this manuscript addresses.

Honest novelty (one paragraph). Relative to PFIb→H3 aggregation papers, we do not improve or reproduce a proprietary building index; we ask whether public flood indicators can support spatially blocked learning on H3. Relative to hexagonal DGGS flood mapping, we emphasise evaluation protocol and adaptive screening rather than scenario inundation alone. Relative to urban pluvial ML susceptibility maps, we add H3-block CV as the primary skill report, an open-label scale-loss ladder, adaptive cell-count diagnostics, and a binding non-PFIb definition of $PFI_h(c, r)$ — while refusing citywide, radar, and rainfall-discrimination claims until stronger evidence exists.

3 Study area and data

Extent. Lower Manhattan bounding box (approx. -74.02 – -73.97°E , 40.70 – 40.76°N), as documented in the public repository configuration and download manifest. This is a **pilot extent**, not citywide New York City.

Live layers (open downloads, August 2026). DEM (USGS 3DEP subset), DEP stormwater flood polygons, building footprints, USGS Ida high-water marks, 311 flooding points (ArcGIS/CDN mirrors), FEMA Sandy inundation (negative control only), NLCD impervious fraction, NHDPlus HR hydro vectors (`dist_stream_m` as a **distance-to-water** proxy in a tidal/shoreline setting). FloodNet is stub/opt-in. Event rainfall is a **synthetic constant** Ida-like hook, not radar.

Provenance. Training table assembly reports open-data assembly with observed static features; rainfall may still carry an event-raster provenance tag rather than gauge/radar observation.

4 Methods

4.1 H3 representation

Training uses H3 resolution **R9** (configurable). Scale-loss diagnostics assemble labels at fine resolution **R10** and roll hotspots to parents R9/R8. Adaptive refinement expands selected parents to **R11**. Joins use EPSG:4326.

4.2 Features and open labels

Static predictors: elevation, slope, flow-accumulation proxy, impervious fraction, urban land-cover flag, building density, distance-to-water. Rainfall enters as rainfall intensity for scenario conditioning; observed static columns are tracked separately from rainfall so a constant synthetic grid is not marketed as radar.

Labels combine DEP polygon area fractions with 311 / Ida point counts into flood-risk / flood-class targets as **observed flood labels** (not ground truth). We do not use PFIb. FloodNet joins only when explicitly enabled and a non-empty layer exists.

4.3 Models and baselines

Primary learner: gradient-boosting classifier + continuous risk regressor. Baselines: L2 logistic / linear, and an elevation–impervious–slope–TWI-like ponding rule. In-sample metrics are optimistic references only.

4.4 Spatial H3-block cross-validation

Cells are grouped by coarse H3 parents; GroupKFold holds out entire blocks. Per-fold metrics are archived with the pilot run. Random i.i.d. splits are not primary.

4.5 Scale-loss diagnostics

Fine-resolution hotspot sets (top quantile of risk) are compared to parent-aggregated hotspots via Jaccard and F1 for mean/max/p90 rollups. Numeric values are **not** equated to Svelling et al.'s PFIb Jaccard ≈ 0.14 .

4.6 Adaptive refinement

After training, cell scores (PFI_h / flood probability) screen parents for refinement. Metrics: mixed-resolution cell counts versus fixed coarse and versus uniform fine (adaptive/uniform cell-count ratio). The pilot uses trained PFI_h as the screening score.

4.7 Event-conditioned PFI_h(c,r)

\[

$$\mathrm{PFI}_h(c,r)=\widehat{P}(Y_c=1\mid X_c,r)$$

\]

Static features (X_c) are held fixed while rainfall condition (r) varies across named scenarios. This is a **model output**, not SHAP/permutation importance and not PFIb.

4.8 Negative control

FEMA Sandy coastal inundation overlaps are reported for separation checks only and are never training labels.

5 Experiments

ID	Description	Role
E1	Assemble open LM table	Feature/label table for the pilot bbox
E2	Spatial CV	Primary blocked accuracy / F1
E3	Baselines	Diagnostic comparison
E4	Jaccard ladder	Scale-loss hotspot similarity
E5	Adaptive + ablation	Cell-count efficiency vs uniform fine
E6	Rainfall scenarios	Within-cell PFI_h response across intensities
E7	Sandy negative control	Coastal overlap checks (not training)

Reproducible tables and figures are released with the public repository; process detail lives in the companion research report.

6 Results

All numbers below are copied from the live Lower Manhattan open-data pilot (n_cells=141). They do **not** describe citywide performance.

6.1 Spatial H3-block CV (primary)

Metric	Value
n_cells	141

spatial_cv_n_folds	5
spatial_cv_n_blocks	7
spatial_cv_accuracy_mean $\pm$ std	0.784 $\pm$ 0.069
spatial_cv_f1_mean	0.866
spatial_cv_r2_mean $\pm$ std	0.030 $\pm$ 0.343
spatial_cv_mae_mean	0.332
random_split_val_accuracy (diagnostic only)	0.690

Per-fold accuracy/F1: Fold0 0.755 / 0.850; Fold1 0.760 / 0.850; Fold2 0.773 / 0.872; Fold3 0.714 / 0.813; Fold4 0.917 / 0.944.

Reading. Mean blocked classification accuracy near 0.78 indicates that an open-label H3 table **can be fit and evaluated** under spatial holdout on this pilot bbox; it does **not** establish operational citywide skill. Fold4 accuracy (0.917) coincides with a small test set ($n_{\text{test}}=24$, two blocks) and must not be cherry-picked. Continuous-risk R^2 (0.030 ± 0.343) is weak and is reported only for completeness. Random-split accuracy (0.690) remains diagnostic only and is **not** the primary claim.

6.2 Scale-loss Jaccard ladder (open labels)

Fine resolution R10, hotspot quantile 0.9:

coarse_res	aggregation	jaccard	f1
8	mean	0.167	0.286
8	max	1.000	1.000
8	p90	1.000	1.000
9	mean	0.977	0.988
9	max	1.000	1.000
9	p90	0.977	0.988

Reading. Under mean aggregation, parent rollup at R8 yields Jaccard **0.167** with fine R10 hotspot parents, which is **consistent with** strong scale smoothing on this open-label stack. Max/p90 retain extrema by construction (Jaccard 1.0 at R8) and therefore should not be narrated as “no scale loss.” These diagnostics are **not** a reproduction of Svelling et al.’s PFIB Jaccard ≈ 0.14 (different labels, resolutions, and hotspot definitions).

6.3 Adaptive vs fixed / uniform fine

Field	Value
-------	-------

score used for screening	PFI_h
n_fixed_coarse (R9)	141
n_adaptive_mixed	3933
n_uniform_fine (R11)	6909
adaptive_cell_count_ratio (adaptive/uniform)	0.569
n_parents_refined	79
score_quantile	0.8

Adaptive mixed grids use fewer cells than uniform R11 (~57%) while refining 79/141 coarse parents. This statement concerns **cell counts only**. We do **not** report wall-clock runtime, memory, or city-scale cost savings from this pilot.

6.4 Rainfall scenarios (PFI_h)

141 cells $\times$ scenarios {moderate 25, heavy 40, ida_like 75, extreme 100 mm/h}; rainfall provenance remains event-raster (synthetic constant hook); assembly is open-data.

Observed: mean PFI_h $\approx$ **0.802888** for every scenario; **within-cell range across scenarios = 0**.

Interpretation: the current pilot does **not** demonstrate rainfall-conditioned discrimination. The scenario loop is correct (only rainfall_mm_h varies); the flat response is explained by **constant training rainfall** — all 141 training cells carry rainfall_mm_h = 75 from the synthetic constant event-raster hook, so rainfall has zero training variance and a classifier feature importance of **0.0**. The model has never observed rainfall variation, so PFI_h cannot respond to it. Non-zero response requires ingested observed event rainfall (multi-intensity, non-synthetic provenance) and retraining; definition of PFI_h(c,r) remains binding for future runs.

6.5 Sandy negative control

Open-data assembly: n_cells=141; n_coastal=31; n_pluvial=71; n_both=23; n_coastal_only=8; frac_coastal_only $\approx$ 0.057; pluvial_minus_coastal_mean_score $\approx$ 0.120. Coastal overlay is **not** a training label.

7 Discussion

7.1 Implications under stated boundaries

The LM pilot **indicates** that an open-label H3 table can be trained and evaluated with blocked CV, that mean rollups can erase fine hotspots (scale-loss), and that trained-score adaptive refinement reduces uniform-fine **cell count**. These results support **protocol credibility** under stated boundaries, not product-ready city maps and not rainfall-responsive screening under the current synthetic rainfall hook.

Relative to Svelling et al. (2026), the dialogue is conceptual: they aggregate a proprietary building index (PFIb) into H3 for scalable communication; we learn on public labels with spatial holdouts and an explicit non-PFIb $PFI_h(c, r)$. Our open-label Jaccard mean R10→R8 (0.167) should **not** be narrated as matching their PFIb Jaccard ≈ 0.14 — different labels, resolutions, and hotspot definitions. Proximity of 0.167 to 0.14 is coincidental and must not be used as validation.

7.2 Limitations

1. **Spatial extent.** Results use a Lower Manhattan pilot (n=141 cells), not citywide New York City.
2. **Label bias.** 311 and related open indicators reflect reporting and mapping processes, not complete ground truth.
3. **Hydro proxy.** Distance-to-water from NHDPlus in a tidal/shoreline setting is a proxy, not inland drainage density.
4. **Rainfall provenance.** Event rainfall remains a synthetic constant hook; gauge/radar event ingest is 待补充.
5. **Flat $PFI_h(c, r)$.** Within-cell range across rainfall scenarios is currently 0; rainfall-conditioned discrimination is not demonstrated.
6. **Continuous skill.** Spatial-CV R^2 is near zero and is not claimed as predictive skill.
7. **Held-out sensors.** FloodNet validation is unavailable under the default stub/opt-in path.
8. **Figure 1.** Workflow schematic is 待补充.

Random-split accuracy must not displace spatial CV in claims. Fold-level variance (including Fold4 on $n_{\text{test}}=24$) further warns against over-interpreting a single pilot run.

7.3 Completion criteria

(i) ingest observed event rainfall with non-synthetic provenance; (ii) non-flat scenarios — within-cell PFI_h range > 0 across rainfall intensities on the same static $\backslash(X_c\backslash)$; (iii) expanded bbox / citywide profile with the same spatial CV protocol; (iv) FloodNet held-out validation when a non-empty layer exists.

8 Conclusions

We present a reproducible open-label H3+ML **protocol** — spatial block CV, scale-loss diagnostics, adaptive cell-count refinement, and an explicit non-PFIb $PFI_h(c, r)$ definition — evaluated on a Lower Manhattan open-data pilot. Live metrics support pipeline credibility under the stated boundaries and do not establish citywide operational skill. Observed event rainfall, non-degenerate rainfall response, expanded spatial extent, and FloodNet validation remain 待补充.

Figure captions

Figure 2. Open-label scale-loss across H3 resolutions. Jaccard similarity and F1 between fine-resolution (R10) hotspot parent sets and coarse (R8/R9) hotspot sets under mean, max, and p90 aggregation. Mean aggregation at R8 yields Jaccard ≈ 0.167 , consistent with smoothing of local extremes; max/p90 preserve extrema by construction. Values are diagnostics on the Lower Manhattan open-label pilot and must not be equated to Svelling et al.'s PFIB Jaccard ≈ 0.14 .

Figure 3. Adaptive refinement versus uniform fine grids (cell counts). Fixed coarse (R9), adaptive mixed, and uniform fine (R11) cell counts for the pilot table (adaptive/uniform ≈ 0.569 ; 79 of 141 parents refined). The comparison is restricted to cell counts; it is not a wall-clock or citywide runtime claim.

Figure 4. Spatial H3-block cross-validation fold metrics. Per-fold classification accuracy and F1 on the Lower Manhattan pilot (n=141 cells; 5 folds). Primary reporting uses the mean $\pm$ standard deviation across folds; individual high folds (e.g. Fold4) are not interpreted as citywide skill.

Figure 1 (workflow schematic). 待补充.

Data and code availability

Public repository (code, configs, tests, paper docs, and small summary tables; large rasters/geojson, trained model binaries, and large parquet files excluded): <https://github.com/Coucou2016/pluvial-flood-risk-DGGS-H3> (pluvial-flood-risk-dggs-h3 v0.1.0). Layer provenance and download manifests are released with the repository. Demo/Oslo paths are QA only. A deep process-oriented research report accompanies the manuscript in the same documentation folder.

References (selected)

1. Svelling, W., Torgersen, G., Bruland, O. & Muthanna, T. Scalable pluvial flood risk assessment: A data-driven framework integrating machine learning (ML) and discrete global grid systems (DGGS H3). **Int. J. Disaster Risk Reduction** **137** (2026). <https://doi.org/10.1016/j.ijdr.2026.106091>
2. Svelling, W. et al. Indexing areas vulnerable to pluvial floods—Using Machine Learning and H3 Hexagonal Grid System. **SSRN** (preprint). <https://doi.org/10.2139/ssrn.5875380>
3. Li, M., McGrath, H. & Stefanakis, E. Multi-Scale Flood Mapping under Climate Change Scenarios in Hexagonal Discrete Global Grids. **ISPRS Int. J. Geo-Inf.** **11**, 627 (2022). <https://doi.org/10.3390/ijgi11120627>
4. Sun, K., Hu, Y., Lakhanpal, G. & Zhou, R. Z. Spatial cross-validation for GeoAI. In Gao, S., Hu, Y. & Li, W. (eds) **Handbook of Geospatial Artificial Intelligence** (Taylor & Francis, 2023). Chapter PDF: https://www.acsu.buffalo.edu/~yhu42/papers/2023_GeoAIHandbook_SpatialCV.pdf
5. Bersabe, J. T. & Jun, B.-W. The Machine Learning-Based Mapping of Urban Pluvial Flood Susceptibility in Seoul Integrating Flood Conditioning Factors and Drainage-Related Data. **ISPRS Int. J. Geo-Inf.** **14**, 57 (2025). <https://doi.org/10.3390/ijgi14020057>

Additional venue-specific references 待补充.